



Hidden-state rank does not reliably transport to out-of-seed pricing performance

● K=1 ■ K=2 ▲ K=3 ◆ K=4 + K=5 × K=8 ○ K-only prediction — geometry improves error — geometry worsens error

A Primary endpoint: HJ loss



B Secondary endpoint: factor Sharpe



Each arrow is one frozen network. Open circle: prediction from nominal K. Colored endpoint: prediction after adding hidden-3 participation rank.